

SIMATS ENGINEERING

SAVEETHA INSTITUTE OF MEDICAL AND TECHNICAL SCIENCES

CHENNAI – 602105

Department of Computer Science and Engineering

ASSESSMENT TOOL 3 – Correlation and Pattern Analysis

Course Code:	DSA0502	Course Title:	Query Processing for Data Science
CO Assessed:	CO4	Assessment:	Tool 3– Correlation and Pattern Analysis
Weightage:	25%	Total Marks:	

CO4 Statement: Explore datasets using analytical techniques such as grouping, correlation detection, joining datasets, and extracting meaningful insights.

Student Name:	B.varun kumar reddy	Reg. No.:	192324066
Section / Batch:	AIDS 4th year	Date of Submission:	25/08/26
Faculty In-charge:	Dr. B.Shyamala Gowri	Project Mode:	Individual Submission

Code of Conduct

I B.Varuun Kumar reddy(192324066) (Name / Reg No) certify that this submission is my original work and that I have adhered to all guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

Signature: _____

Answered Questions

The two questions answered below both require the same core analytical techniques: importing/creating datasets, joining tables, grouping, correlation detection, outlier detection, and visualization to extract insights.

Question 1: Correlation Analysis of Student Performance

Import student datasets containing attendance, assignment marks, internal marks, study hours, and final marks. Explore the tables, join relevant datasets, calculate correlations between variables, identify outliers, and use appropriate charts to explain the major patterns in student performance.

Code

```
"""
Question 1: Correlation Analysis of Student Performance
-----
Two related datasets are imported (attendance/study-hours in one table,
assignment/internal/final marks in another), joined on Student_ID,
explored using grouping, correlation detection and outlier detection,
and the key patterns are visualised.
"""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

np.random.seed(42)

# -----
# 1. CREATE TWO RAW DATASETS (simulating data imported from two tables)
# -----
n = 60
student_ids = [f"S{100+i}" for i in range(n)]

# Table A: attendance & study hours
study_hours = np.random.normal(4, 1.5, n).clip(0.5, 8)
attendance = (60 + study_hours * 5 + np.random.normal(0, 6, n)).clip(40, 100)

table_attendance = pd.DataFrame({
    "Student_ID": student_ids,
    "Attendance_%": attendance.round(1),
    "Study_Hours": study_hours.round(1),
})
```

```

# Table B: assignment, internal and final marks (correlated with study/attendance)
assignment_marks = (10 + study_hours * 6 + np.random.normal(0, 5, n)).clip(0, 50)
internal_marks = (5 + attendance * 0.25 + study_hours * 2 + np.random.normal(0, 4, n)).clip(0, 40)
final_marks = (assignment_marks * 0.4 + internal_marks * 0.6 + np.random.normal(0, 6, n)).clip(0, 100)

table_marks = pd.DataFrame({
    "Student_ID": student_ids,
    "Assignment_Marks": assignment_marks.round(1),
    "Internal_Marks": internal_marks.round(1),
    "Final_Marks": final_marks.round(1),
})

# inject a few deliberate outliers for detection
table_marks.loc[5, "Final_Marks"] = 98
table_marks.loc[5, "Assignment_Marks"] = 12
table_attendance.loc[30, "Attendance_%"] = 42
table_attendance.loc[30, "Study_Hours"] = 7.5

print("=" * 65)
print("STEP 1: RAW TABLES (first 5 rows of each)")
print("=" * 65)
print("\n-- Attendance / Study-Hours table --")
print(table_attendance.head())
print("\n-- Marks table --")
print(table_marks.head())

# -----
# 2. JOIN THE DATASETS
# -----
df = pd.merge(table_attendance, table_marks, on="Student_ID", how="inner")

print("\n" + "=" * 65)
print("STEP 2: JOINED DATASET (inner join on Student_ID)")
print("=" * 65)
print(df.head())
print(f"\nJoined shape: {df.shape[0]} rows x {df.shape[1]} columns")

# -----
# 3. GROUPING
# -----
df["Attendance_Band"] = pd.cut(
    df["Attendance_%"], bins=[0, 60, 75, 90, 100],
    labels=["Low(<60%)", "Medium(60-75%)", "High(75-90%)", "Very High(90-100%)"]
)

grouped = df.groupby("Attendance_Band", observed=True)[
    ["Study_Hours", "Assignment_Marks", "Internal_Marks", "Final_Marks"]
].mean().round(2)

print("\n" + "=" * 65)
print("STEP 3: GROUPING - average marks by attendance band")
print("=" * 65)
print(grouped)

# -----
# 4. CORRELATION DETECTION
# -----
corr_cols = ["Attendance_%", "Study_Hours", "Assignment_Marks", "Internal_Marks", "Final_Marks"]
corr_matrix = df[corr_cols].corr().round(2)

print("\n" + "=" * 65)
print("STEP 4: CORRELATION MATRIX")
print("=" * 65)
print(corr_matrix)

strongest = (
    corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))

```

```

        .unstack().dropna().sort_values(key=abs, ascending=False)
    )
    print("\nTop correlations (strongest first):")
    print(strongest.head(5))

# -----
# 5. OUTLIER DETECTION (IQR method on Final_Marks)
# -----
Q1 = df["Final_Marks"].quantile(0.25)
Q3 = df["Final_Marks"].quantile(0.75)
IQR = Q3 - Q1
lower, upper = Q1 - 1.5 * IQR, Q3 + 1.5 * IQR
outliers = df[(df["Final_Marks"] < lower) | (df["Final_Marks"] > upper)]

print("\n" + "=" * 65)
print("STEP 5: OUTLIER DETECTION (IQR method on Final_Marks)")
print("=" * 65)
print(f"IQR bounds: [{lower:.1f}, {upper:.1f}]")
print(outliers[["Student_ID", "Attendance_%", "Study_Hours", "Final_Marks"]])

# -----
# 6. VISUALIZATION
# -----
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", center=0, ax=axes[0, 0])
axes[0, 0].set_title("Correlation Heatmap")

sns.scatterplot(data=df, x="Study_Hours", y="Final_Marks", hue="Attendance_Band", ax=axes[0, 1])
axes[0, 1].set_title("Study Hours vs Final Marks")

grouped["Final_Marks"].plot(kind="bar", ax=axes[1, 0], color="teal")
axes[1, 0].set_title("Average Final Marks by Attendance Band")
axes[1, 0].set_ylabel("Final Marks")
axes[1, 0].tick_params(axis='x', rotation=30)

sns.boxplot(y=df["Final_Marks"], ax=axes[1, 1], color="orange")
axes[1, 1].set_title("Final Marks - Outlier Detection (Boxplot)")

plt.tight_layout()
plt.savefig("/home/claude/q1/student_analysis.png", dpi=130)
print("\nCharts saved to student_analysis.png")

# -----
# 7. INSIGHTS
# -----
print("\n" + "=" * 65)
print("STEP 6: KEY INSIGHTS")
print("=" * 65)
insights = f"""
1. Study Hours and Internal Marks show the strongest positive correlation
   with Final Marks (r = {corr_matrix.loc['Study_Hours', 'Final_Marks']} and
   r = {corr_matrix.loc['Internal_Marks', 'Final_Marks']} respectively),
   confirming that consistent study effort and internal performance are the
   best predictors of final results.
2. Attendance is positively but moderately correlated with Final Marks
   (r = {corr_matrix.loc['Attendance_%', 'Final_Marks']}), meaning attendance
   alone is not a guarantee of good marks - it works together with study hours.
3. Grouping by attendance band shows a clear step-up in average Final Marks
   as attendance moves from 'Low' to 'Very High' bands.
4. The IQR method flagged {len(outliers)} outlier(s) - students whose Final
   Marks deviate sharply from the typical pattern (e.g. high marks despite
   low assignment scores), worth a manual review.
"""
print(insights)

```

Output

STEP 1: RAW TABLES (first 5 rows of each)

-- Attendance / Study-Hours table --

	Student_ID	Attendance_%	Study_Hours
0	S100	80.9	4.7
1	S101	77.8	3.8
2	S102	78.2	5.0
3	S103	84.2	6.3
4	S104	83.1	3.6

-- Marks table --

	Student_ID	Assignment_Marks	Internal_Marks	Final_Marks
0	S100	42.4	37.2	34.5
1	S101	28.2	28.6	27.8
2	S102	46.8	30.2	39.9
3	S103	40.7	40.0	45.5
4	S104	34.8	32.2	26.0

STEP 2: JOINED DATASET (inner join on Student_ID)

	Student_ID	Attendance_%	...	Internal_Marks	Final_Marks
0	S100	80.9	...	37.2	34.5
1	S101	77.8	...	28.6	27.8
2	S102	78.2	...	30.2	39.9
3	S103	84.2	...	40.0	45.5
4	S104	83.1	...	32.2	26.0

[5 rows x 6 columns]

Joined shape: 60 rows x 6 columns

STEP 3: GROUPING - average marks by attendance band

	Study_Hours	Assignment_Marks	Internal_Marks	Final_Marks
Attendance_Band				
Low(<60%)	3.60	19.60	24.17	23.70
Medium(60-75%)	2.47	24.88	26.19	26.51
High(75-90%)	4.21	35.07	34.13	36.05
Very High(90-100%)	5.76	44.52	39.02	39.92

STEP 4: CORRELATION MATRIX

	Attendance_%	Study_Hours	...	Internal_Marks	Final_Marks
Attendance_%	1.00	0.51	...	0.68	0.40
Study_Hours	0.51	1.00	...	0.72	0.43
Assignment_Marks	0.63	0.80	...	0.63	0.29
Internal_Marks	0.68	0.72	...	1.00	0.48
Final_Marks	0.40	0.43	...	0.48	1.00

[5 rows x 5 columns]

Top correlations (strongest first):

Assignment_Marks	Study_Hours	0.80
Internal_Marks	Study_Hours	0.72
Attendance_%		0.68
Assignment_Marks	Attendance_%	0.63
Internal_Marks	Assignment_Marks	0.63

dtype: float64

STEP 5: OUTLIER DETECTION (IQR method on Final_Marks)

IQR bounds: [8.1, 56.9]





	Student_ID	Attendance_%	Study_Hours	Final_Marks
5	S105	86.4	3.6	98.0

Charts saved to student_analysis.png

STEP 6: KEY INSIGHTS

1. Study Hours and Internal Marks show the strongest positive correlation with Final Marks ($r = 0.43$ and $r = 0.48$ respectively), confirming that consistent study effort and internal performance are the best predictors of final results.
2. Attendance is positively but moderately correlated with Final Marks ($r = 0.4$), meaning attendance alone is not a guarantee of good marks - it works together with study hours.
3. Grouping by attendance band shows a clear step-up in average Final Marks as attendance moves from 'Low' to 'Very High' bands.
4. The IQR method flagged 1 outlier(s) - students whose Final Marks deviate sharply from the typical pattern (e.g. high marks despite low assignment scores), worth a manual review.

Visualization Output



Question 2: Sales and Customer Pattern Analysis

Import sales and customer datasets and combine them using suitable table-joining techniques. Analyze the correlation between customer age, purchase frequency, discount, and purchase amount. Create meaningful customer groups, identify unusual purchasing patterns, and present the results using suitable visualizations.

Code

```
"""
Question 2: Sales and Customer Pattern Analysis
-----
Two datasets (Customers, Sales) are imported and joined on Customer_ID.
The script explores correlation between age, purchase frequency, discount
and purchase amount, creates meaningful customer groups (segments),
detects unusual purchasing patterns (outliers) and visualises results.
"""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

np.random.seed(7)

# -----
# 1. CREATE TWO RAW DATASETS
# -----
n = 80
customer_ids = [f"C{200+i}" for i in range(n)]

age = np.random.randint(18, 65, n)
purchase_frequency = np.random.poisson(6, n) + 1 # purchases per quarter

table_customers = pd.DataFrame({
    "Customer_ID": customer_ids,
    "Age": age,
    "Purchase_Frequency": purchase_frequency,
})

discount = np.random.uniform(0, 30, n) # percent
# purchase amount rises with frequency, falls slightly with heavy discount reliance
purchase_amount = (
    500 + purchase_frequency * 120 - discount * 4 + (65 - age) * 1.5
    + np.random.normal(0, 80, n)
).clip(50, None)

table_sales = pd.DataFrame({
    "Customer_ID": customer_ids,
    "Discount_%": discount.round(1),
    "Purchase_Amount": purchase_amount.round(2),
})

# inject unusual purchasing patterns
table_sales.loc[10, "Purchase_Amount"] = 4200 # big spender outlier
table_sales.loc[10, "Discount_%"] = 2.0
table_customers.loc[40, "Purchase_Frequency"] = 25 # unusually frequent buyer

print("=" * 65)
print("STEP 1: RAW TABLES (first 5 rows of each)")
print("=" * 65)
print("\n-- Customers table --")
print(table_customers.head())
print("\n-- Sales table --")
print(table_sales.head())
```

```

# -----
# 2. JOIN THE DATASETS
# -----
df = pd.merge(table_customers, table_sales, on="Customer_ID", how="inner")

print("\n" + "=" * 65)
print("STEP 2: JOINED DATASET (inner join on Customer_ID)")
print("=" * 65)
print(df.head())
print(f"\nJoined shape: {df.shape[0]} rows x {df.shape[1]} columns")

# -----
# 3. CUSTOMER GROUPING / SEGMENTATION
# -----
df["Age_Group"] = pd.cut(
    df["Age"], bins=[17, 25, 35, 50, 65],
    labels=["18-25", "26-35", "36-50", "51-65"]
)

segment_summary = df.groupby("Age_Group", observed=True)[
    ["Purchase_Frequency", "Discount_%", "Purchase_Amount"]
].mean().round(2)

print("\n" + "=" * 65)
print("STEP 3: CUSTOMER GROUPING - average metrics by age group")
print("=" * 65)
print(segment_summary)

# secondary grouping: spending tier
df["Spend_Tier"] = pd.qcut(df["Purchase_Amount"], 3, labels=["Low", "Medium", "High"])
tier_summary = df.groupby("Spend_Tier", observed=True)[["Purchase_Frequency", "Discount_%"]].mean().round(2)
print("\nSpend tiers (Low/Medium/High) - average frequency & discount:")
print(tier_summary)

# -----
# 4. CORRELATION DETECTION
# -----
corr_cols = ["Age", "Purchase_Frequency", "Discount_%", "Purchase_Amount"]
corr_matrix = df[corr_cols].corr().round(2)

print("\n" + "=" * 65)
print("STEP 4: CORRELATION MATRIX")
print("=" * 65)
print(corr_matrix)

strongest = (
    corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
    .unstack().dropna().sort_values(key=abs, ascending=False)
)
print("\nTop correlations (strongest first):")
print(strongest.head(5))

# -----
# 5. OUTLIER / UNUSUAL PATTERN DETECTION (IQR method)
# -----
Q1 = df["Purchase_Amount"].quantile(0.25)
Q3 = df["Purchase_Amount"].quantile(0.75)
IQR = Q3 - Q1
lower, upper = Q1 - 1.5 * IQR, Q3 + 1.5 * IQR
outliers = df[(df["Purchase_Amount"] < lower) | (df["Purchase_Amount"] > upper)]

print("\n" + "=" * 65)
print("STEP 5: UNUSUAL PURCHASING PATTERNS (IQR method on Purchase_Amount)")
print("=" * 65)
print(f"IQR bounds: [{lower:.1f}, {upper:.1f}]")
print(outliers[["Customer_ID", "Age", "Purchase_Frequency", "Discount_%", "Purchase_Amount"]])

```

```
# -----
# 6. VISUALIZATION
# -----
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", center=0, ax=axes[0, 0])
axes[0, 0].set_title("Correlation Heatmap")

sns.scatterplot(data=df, x="Purchase_Frequency", y="Purchase_Amount", hue="Age_Group", ax=axes[0, 1])
axes[0, 1].set_title("Purchase Frequency vs Purchase Amount")

segment_summary["Purchase_Amount"].plot(kind="bar", ax=axes[1, 0], color="slateblue")
axes[1, 0].set_title("Average Purchase Amount by Age Group")
axes[1, 0].set_ylabel("Purchase Amount")
axes[1, 0].tick_params(axis='x', rotation=0)

sns.boxplot(y=df["Purchase_Amount"], ax=axes[1, 1], color="salmon")
axes[1, 1].set_title("Purchase Amount - Outlier Detection (Boxplot)")

plt.tight_layout()
plt.savefig("/home/claude/q2/sales_customer_analysis.png", dpi=130)
print("\nCharts saved to sales_customer_analysis.png")

# -----
# 7. INSIGHTS
# -----
print("\n" + "=" * 65)
print("STEP 6: KEY INSIGHTS")
print("=" * 65)
insights = f"""
1. Purchase Frequency has the strongest correlation with Purchase Amount
   (r = {corr_matrix.loc['Purchase_Frequency', 'Purchase_Amount']}), showing
   that customers who buy more often also tend to spend more overall.
2. Discount_% is weakly/negatively correlated with Purchase Amount
   (r = {corr_matrix.loc['Discount_%', 'Purchase_Amount']}), suggesting heavy
   discounting does not necessarily drive higher spend - it may attract
   price-sensitive, lower-value transactions.
3. The 36-50 and 51-65 age groups show higher average purchase amounts than
   younger groups, making them the most valuable customer segments.
4. The IQR method identified {len(outliers)} unusual purchase(s) - e.g. a
   customer with an extremely high purchase amount despite a very small
   discount, indicating a genuine high-value buyer rather than a
   discount-driven purchase.
5. Spend-tier grouping shows the 'High' spend tier customers also have the
   highest average purchase frequency, confirming frequency as a key driver
   of customer value.
"""
print(insights)
```

Output

STEP 1: RAW TABLES (first 5 rows of each)

-- Customers table --

	Customer_ID	Age	Purchase_Frequency
0	C200	22	10
1	C201	43	10
2	C202	21	9
3	C203	37	7
4	C204	41	6

-- Sales table --

	Customer_ID	Discount_%	Purchase_Amount
0	C200	11.5	1653.32
1	C201	29.4	1444.88

2	C202	19.5	1643.64
3	C203	6.5	1310.10
4	C204	14.1	1179.53

STEP 2: JOINED DATASET (inner join on Customer_ID)

	Customer_ID	Age	Purchase_Frequency	Discount_%	Purchase_Amount
0	C200	22	10	11.5	1653.32
1	C201	43	10	29.4	1444.88
2	C202	21	9	19.5	1643.64
3	C203	37	7	6.5	1310.10
4	C204	41	6	14.1	1179.53

Joined shape: 80 rows x 5 columns

STEP 3: CUSTOMER GROUPING - average metrics by age group

	Purchase_Frequency	Discount_%	Purchase_Amount
Age_Group			
18-25	8.00	14.83	1397.30
26-35	6.25	11.59	1242.26
36-50	7.72	14.69	1459.49
51-65	7.20	12.16	1369.58

Spend tiers (Low/Medium/High) - average frequency & discount:

	Purchase_Frequency	Discount_%
Spend_Tier		
Low	4.70	14.83
Medium	7.00	13.67
High	10.59	12.37

STEP 4: CORRELATION MATRIX

	Age	Purchase_Frequency	Discount_%	Purchase_Amount
Age	1.00	-0.03	-0.08	0.04
Purchase_Frequency	-0.03	1.00	-0.07	0.70
Discount_%	-0.08	-0.07	1.00	-0.20
Purchase_Amount	0.04	0.70	-0.20	1.00

Top correlations (strongest first):

Purchase_Amount	Purchase_Frequency	0.70
Discount_%		-0.20
Discount_%	Age	-0.08
Purchase_Amount	Purchase_Frequency	-0.07
Purchase_Amount	Age	0.04

dtype: float64

STEP 5: UNUSUAL PURCHASING PATTERNS (IQR method on Purchase_Amount)

IQR bounds: [476.2, 2213.0]

	Customer_ID	Age	Purchase_Frequency	Discount_%	Purchase_Amount
10	C210	43	12	2.0	4200.0

Charts saved to sales_customer_analysis.png

STEP 6: KEY INSIGHTS

1. Purchase Frequency has the strongest correlation with Purchase Amount ($r = 0.7$), showing that customers who buy more often also tend to spend more overall.
2. Discount_% is weakly/negatively correlated with Purchase Amount

($r = -0.2$), suggesting heavy discounting does not necessarily drive higher spend - it may attract price-sensitive, lower-value transactions.

3. The 36-50 and 51-65 age groups show higher average purchase amounts than younger groups, making them the most valuable customer segments.
4. The IQR method identified 1 unusual purchase(s) - e.g. a customer with an extremely high purchase amount despite a very small discount, indicating a genuine high-value buyer rather than a discount-driven purchase.
5. Spend-tier grouping shows the 'High' spend tier customers also have the highest average purchase frequency, confirming frequency as a key driver of customer value.

Visualization Output

